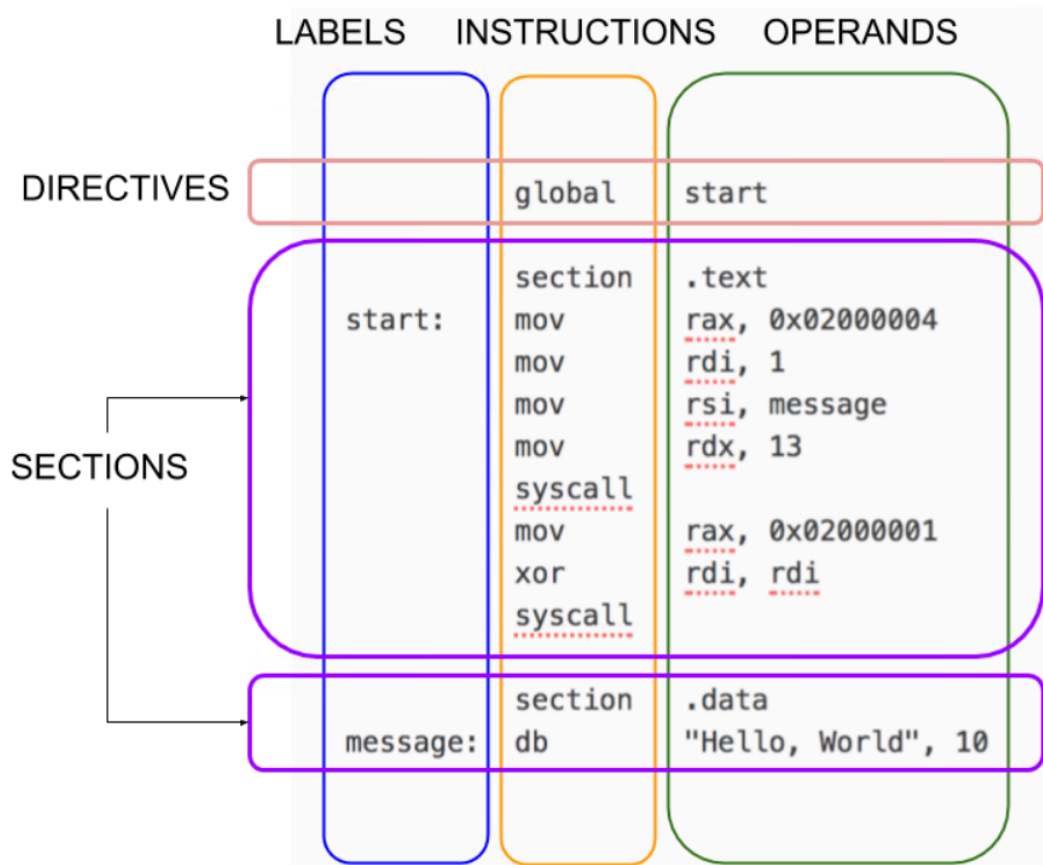




```
; -----  
; Writes "Hello, World" to the console using only system calls. Runs on 64-bit Linux only.  
; To assemble and run:  
;  
;     nasm -felf64 hello.asm && ld hello.o && ./a.out  
; -----  
  
        global  _start  
  
        section .text  
_start:  mov     rax, 1           ; system call for write  
        mov     rdi, 1           ; file handle 1 is stdout  
        mov     rsi, message      ; address of string to output  
        mov     rdx, 13          ; number of bytes  
        syscall                  ; invoke operating system to do the write  
        mov     rax, 60          ; system call for exit  
        xor     rdi, rdi         ; exit code 0  
        syscall                  ; invoke operating system to exit  
  
        section .data  
message: db      "Hello, World", 10 ; note the newline at the end
```

```
$ nasm -felf64 hello.asm && ld hello.o && ./a.out  
Hello, World
```

NASM is line-based. Most programs consist of **directives** followed by one or more **sections**. Lines can have an optional **label**. Most lines have an **instruction** followed by zero or more **operands**.



Generally, you put code in a section called `.text` and your constant data in a section called `.data`.

There are hundreds of instructions. You can't learn them all at once. Just start with these:

<code>mov x, y</code>	$x \leftarrow y$
<code>and x, y</code>	$x \leftarrow x \text{ and } y$
<code>or x, y</code>	$x \leftarrow x \text{ or } y$
<code>xor x, y</code>	$x \leftarrow x \text{ xor } y$
<code>add x, y</code>	$x \leftarrow x + y$
<code>sub x, y</code>	$x \leftarrow x - y$
<code>inc x</code>	$x \leftarrow x + 1$
<code>dec x</code>	$x \leftarrow x - 1$
<code>syscall</code>	Invoke an operating system routine
<code>db</code>	A pseudo-instruction that declares bytes that will be in memory when the program runs

Register Operands

In this tutorial we only care about the integer registers and the xmm registers. You should already know what the registers are, but here is a quick review. The 16 integer registers are 64 bits wide and are called:

R0 R1 R2 R3 R4 R5 R6 R7 R8 R9 R10 R11 R12 R13 R14 R15
RAX RCX RDX RBX RSP RBP RSI RDI

(Note that 8 of the registers have alternate names.) You can treat the lowest 32-bits of each register as a register itself but using these names:

R0D R1D R2D R3D R4D R5D R6D R7D R8D R9D R10D R11D R12D R13D R14D R15D
EAX ECX EDX EBX ESP EBP ESI EDI

You can treat the lowest 16-bits of each register as a register itself but using these names:

R0W R1W R2W R3W R4W R5W R6W R7W R8W R9W R10W R11W R12W R13W R14W R15W
AX CX DX BX SP BP SI DI

You can treat the lowest 8-bits of each register as a register itself but using these names:

R0B R1B R2B R3B R4B R5B R6B R7B R8B R9B R10B R11B R12B R13B R14B R15B
AL CL DL BL SPL BPL SIL DIL


For historical reasons, bits 15 through 8 of R0..R3 are named:

AH CH DH BH

And finally, there are 16 XMM registers, each 128 bits wide, named:

XMM0 ... XMM15

Study this picture; hopefully it helps:



Memory Operands

These are the basic forms of addressing:

- [number]
- [reg]
- [reg + reg*scale] *scale is 1, 2, 4, or 8 only*
- [reg + number]
- [reg + reg*scale + number]

The number is called the **displacement**; the plain register is called the **base**; the register with the scale is called the **index**.

```
[750]           ; displacement only
[rbp]           ; base register only
[rcx + rsi*4]   ; base + index * scale
[rbp + rdx]     ; scale is 1
[rbx - 8]       ; displacement is -8
[rax + rdi*8 + 500] ; all four components
[rbx + counter] ; uses the address of the variable 'counter' as the displacement
```

Immediate Operands

These can be written in many ways. Here are some examples from the official docs.

```
200           ; decimal
0200          ; still decimal - the leading 0 does not make it octal
0200d        ; explicitly decimal - d suffix
0d200        ; also decimal - 0d prefix
0c8h         ; hex - h suffix, but leading 0 is required because c8h looks like a var
0xc8         ; hex - the classic 0x prefix
0hc8         ; hex - for some reason NASM likes 0h
310q         ; octal - q suffix
0q310        ; octal - 0q prefix
11001000b    ; binary - b suffix
0b1100_1000  ; binary - 0b prefix, and by the way, underscores are allowed
```

add *reg, reg*

add *reg, mem*

add *reg, imm*

add *mem, reg*

add *mem, imm*

```

db    0x55                ; just the byte 0x55
db    0x55,0x56,0x57      ; three bytes in succession
db    'a',0x55            ; character constants are OK
db    'hello',13,10,'$'   ; so are string constants
dw    0x1234              ; 0x34 0x12
dw    'a'                 ; 0x61 0x00 (it's just a number)
dw    'ab'                ; 0x61 0x62 (character constant)
dw    'abc'               ; 0x61 0x62 0x63 0x00 (string)
dd    0x12345678          ; 0x78 0x56 0x34 0x12
dd    1.234567e20         ; floating-point constant
dq    0x123456789abcdef0 ; eight byte constant
dq    1.234567e20         ; double-precision float
dt    1.234567e20         ; extended-precision float

```

To reserve space (without initializing), you can use the following pseudo instructions. They should go in a section called `.bss` (you'll get an error if you try to use them in a `.text` section):

```

buffer:      resb    64      ; reserve 64 bytes
wordvar:     resw     1      ; reserve a word
realarray:   resq    10     ; array of ten reals

```

```

        global  start
        section .text

start:

        mov     rdx, output          ; rdx holds address of next byte to write
        mov     r8, 1                ; initial line length
        mov     r9, 0                ; number of stars written on line so far

line:
        mov     byte [rdx], '*'      ; write single star
        inc     rdx                  ; advance pointer to next cell to write
        inc     r9                   ; "count" number so far on line
        cmp     r9, r8               ; did we reach the number of stars for this line?
        jne     line                 ; not yet, keep writing on this line

lineDone:
        mov     byte [rdx], 10       ; write a new line char
        inc     rdx                  ; and move pointer to where next char goes
        inc     r8                   ; next line will be one char longer
        mov     r9, 0                ; reset count of stars written on this line
        cmp     r8, maxlines         ; wait, did we already finish the last line?
        jng     line                 ; if not, begin writing this line

done:
        mov     rax, 0x02000004      ; system call for write
        mov     rdi, 1               ; file handle 1 is stdout
        mov     rsi, output          ; address of string to output
        mov     rdx, dataSize        ; number of bytes
        syscall                       ; invoke operating system to do the write
        mov     rax, 0x02000001      ; system call for exit
        xor     rdi, rdi              ; exit code 0
        syscall                       ; invoke operating system to exit

        section .bss
maxlines equ 8
dataSize equ 44
output:  resb dataSize

```

```
$ nasm -fmacho64 triangle.asm && ld triangle.o && ./a.out
```

```

*
**
***
****
*****
*****
*****
*****
*****

```

- `cmp` does a comparison
- `je` jumps to a label if the previous comparison was equal. We also have `jne` (jump if not equal), `j1` (jump if less), `jnl` (jump if not less), `jg` (jump if greater), `jng` (jump if not greater), `jle` (jump if less or equal), `jnl` (jump if not less or equal), `jge` (jump if greater or equal), `jnge` (jump if not greater or equal), and many more.
- `equ` is actually not a real instruction. It simply defines an abbreviation for the assembler itself to use. (This is a profound idea.)
- The `.bss` section is for *writable* data.

Writing standalone programs with just system calls is cool, but rare. We would like to use the good stuff in the C library.

Remember how in C execution “starts” at the function `main`? That’s because the C library actually has the `_start` label inside itself! The code at `_start` does some initialization, then it calls `main`, then it does some clean up, then it issues the system call for exit. So you just have to implement `main`. We can do that in assembly!

```

global _main
extern _puts

_main:
    section .text
    push    rbx                ; Call stack must be aligned
    lea     rdi, [rel message] ; First argument is address of message
    call    _puts              ; puts(message)
    pop     rbx                ; Fix up stack before returning
    ret

    section .data
message: db "Hola, mundo", 0 ; C strings need a zero byte at the end

```

```

$ nasm -fmacho64 hola.asm && gcc hola.o && ./a.out
Hola, mundo

```